

Linear Regression

$$y = w \cdot x + b$$

$$y = w_1 \cdot x_1 \dots w_n \cdot x_n + b \quad \left. \vphantom{y = w_1 \cdot x_1 \dots w_n \cdot x_n + b} \right\} \text{prediction } y$$

$$\text{Loss} = \frac{1}{m} \sum_{i=1}^m (y_i - t_i)^2$$

sum of all errors \rightarrow squared
m \rightarrow number of examples

$\left. \vphantom{\text{Loss} = \frac{1}{m} \sum_{i=1}^m (y_i - t_i)^2} \right\}$ compare y with t and calculate loss

While loss > precision:

$$w_i = w_i - lr \cdot (y - t) \cdot x_i$$

$\left. \vphantom{w_i = w_i - lr \cdot (y - t) \cdot x_i} \right\}$ update the weights until the loss is close enough to zero

Binary Logistic Regression

$$y = w_1 \cdot x_1 \dots w_n \cdot x_n + b \quad \left. \vphantom{y = w_1 \cdot x_1 \dots w_n \cdot x_n + b} \right\} \text{prediction } y$$

$$\hat{y} = \sigma(y) = \frac{1}{1 + e^{-y}} \quad \left. \vphantom{\hat{y} = \sigma(y) = \frac{1}{1 + e^{-y}}} \right\} \text{apply the sigmoid to } y$$

$\hookrightarrow z =$ probability value between 0 and 1

$$\hat{y}' = \sigma'(y) = \hat{y} \cdot (1 - \hat{y})$$

$\left. \vphantom{\hat{y}' = \sigma'(y) = \hat{y} \cdot (1 - \hat{y})} \right\}$ calculate the derivative of the sigmoid

cross-entropy loss

$$\text{loss} = \frac{1}{m} \sum_{i=1}^m (t_i \cdot \log(\hat{y}_i) + (1-t_i) \cdot \log(1-\hat{y}_i))$$

* calculate the gradients of the loss function

by taking the partial derivatives of the loss function with respect to each parameter

While loss > precision:

$$w_i = w_i - \text{lr} (y - t) x_i$$

} update the weights until the loss is close enough to zero

Multinomial Logistic Regression

$$y = w_1 \cdot x_1 \dots w_n \cdot x_n + b \quad \text{prediction } y$$

$$\hat{y} = \frac{e^{y_j}}{\sum_{k=1}^n e^{y_k}} \quad \left\{ \begin{array}{l} \text{take the exponential of the} \\ \text{weighted sum and normalize} \\ \text{divide each exponential of} \\ \text{the weighted sum by the sum} \\ \text{of all the exponentials} \end{array} \right. \rightarrow \text{softmax}$$

cross-entropy loss

$$\text{loss} = \frac{1}{m} \sum_{i=1}^m (t_i \cdot \log(\hat{y}_i) + (1-t_i) \cdot \log(1-\hat{y}_i))$$

While loss > precision:

$$w_i = w_i - \text{lr} (y - t) x_i$$

} update the weights until the loss is close enough to zero

Perceptron

for each neuron j :

$$y_j = w_1 \cdot x_1 \dots w_n \cdot x_n$$

$$y_j = \begin{cases} 1 & \text{if } \sum_{i=0}^m w_{ij} \cdot x_i > 0 \\ 0 & \text{if } \sum_{i=0}^m w_{ij} \cdot x_i \leq 0 \end{cases}$$

compute activation function

$$w_{i,j} = w_{i,j} - \text{lr} (y_j - t_j) x_i \quad \left. \vphantom{w_{i,j}} \right\} \text{update weights individually}$$

Neural Networks

- ① Multiply input matrix X with a weight matrix W , and add a bias

$$y = w_1 \cdot x_1 \dots w_n \cdot x_n + b = \sum_{i=0}^m w_{ij} \cdot x_i$$

- ② Use a differentiable activation function (element-wise)

$$\hat{y} = \sigma(y) = \frac{1}{1 + e^{-y}} \quad \text{OR} \quad \hat{y} = \max(y, 0)$$

↳ Returns a new matrix

- ③ Repeat step ① and ② for each layer

④ At the final layer, use the sigmoid for binary neural networks and multi-label neural networks and the softmax for multi-class neural networks.

⑤ Apply a loss function \rightarrow MSE or Cross-entropy L_o to the output of the activation function from the final layer (sigmoid or softmax)

⑥ Compute the delta terms in the output layer and update the weights between the output layer and hidden layer

$\delta_o(k_j)$ = delta term at the final node j
 $\delta_o(k_j) = y - t$ (assumes cross-entropy loss)

⑦ Compute the delta terms in the hidden layer

$\delta(\text{hidden}_j) = \underbrace{\tilde{a}_j(1 - a_j)}_{\text{output after activation of node } j} \cdot \sum_{i=1}^n \underbrace{\delta_o(k_i)}_{\text{delta term at output node } i} \cdot \underbrace{w_{j,i}}_{\text{Weight from hidden node } j \text{ to output node } i}$
for $j = 1, \dots, k$

* For one instance

$$\delta(h_j) = (y - t) \cdot a_j(1 - a_j) \cdot w_{ji}$$

⑧ Update the weights by the deltas in both layers

$$w_{i,j} = w_{i,j} - lr \cdot \delta_o(k_j) \cdot a_j$$

$$v_{i,j} = v_{i,j} - lr \cdot \delta(\text{hidden}_j) \cdot x_i$$

Scaling, Regularization and Euclidean Distance

min-max-scaling \rightarrow for each feature, find the minimum and maximum value and transform the feature to numbers in the interval 0,1.

$$\frac{x - \min}{\max - \min}$$

Normalizer

$$\frac{x - \text{mean}}{\text{std}}$$


Regularization \rightarrow prevents overfitting

\hookrightarrow adds a penalty term to the model's loss function

$$\text{loss}(X, t, w) = a \cdot R(w)$$

$$R(w) = \vec{w} \cdot \vec{w}$$

Euclidean distance \rightarrow measure of the straight-line distance between two points in a multi-dimensional space.



Reinforcement Learning

greedy action selection

$\hookrightarrow \max Q(s, a) \rightarrow$ action with the highest average reward value in the Q -table for state s .

ϵ -greedy action selection:

\hookrightarrow initialize a probability $\epsilon \in [0, 1]$
select $\max Q(s, a)$ with prob $1 - \epsilon$
select random a with prob ϵ

Update rules:

* SARSA *

reward for current action

new state value (assuming current policy)

discount rate

current value

learning rate

Q-table value

$$Q[s, a] = (1 - \mu) \cdot Q[s, a] + \mu \cdot (r + \gamma \cdot Q[s', a'])$$

* Q-learning *

new state value, assuming a greedy policy

$$Q[s, a] = (1 - \mu) \cdot Q[s, a] + \mu \cdot (r + \gamma \cdot \max_{a'} Q[s', a'])$$

Distance measures

Manhattan

$$d(x_1, x_2) = \sum_{i=1}^p |x_{1i} - x_{2i}|$$

Two datapoints

absolute difference between each datapoint

sum of all differences

Euclidean

$$d(x_1, x_2) = \sqrt{\sum_{i=1}^p (x_{1i} - x_{2i})^2}$$

(1) difference
 (2) square the differences
 (3) sum
 (4) take the square root

Cosine

$$d(x_1, x_2) = 1 - \frac{x_1 \cdot x_2}{\|x_1\| \cdot \|x_2\|}$$

cos θ
 length of vector

EVALUATION

True Positive

False Positive

False Negative

True Negative

Recall

$$\frac{TP}{TP + FN}$$

Precision

$$\frac{TP}{TP + FP}$$

Accuracy

$$\frac{TP + TN}{TP + FP + TN + FN}$$

F-Score →

$$\frac{2 \cdot P \cdot R}{P + R}$$