

# **IN5170 Models of Concurrency**

## Lecture 1: Introduction

---

Andrea Pferscher

August 20th, 2025

University of Oslo

# Motivation

Concurrency is everywhere and challenging

- Multiple computations at the same time
- Distributed systems, cloud computing, networks, programs, . . . , even single-threaded application!
- Source of serious flaws in systems: *safety, security, privacy*
- Hard or impossible to test concurrent systems – enormous amount of executions

Concurrency has many facets

- Programming languages use many different concurrency mechanisms
  - Multi-threading with shared state
  - Message passing, channels
  - Go-routines, async/await, futures
  - ...
- What is the essence of these mechanisms and how should we use them?

## Learning outcomes

- Knowledge about different mechanisms of parallelism, including shared variables and communication-based techniques
- Good insight into typical problems with parallel systems like deadlock, fairness, etc.
- Specification, design and analysis of a parallel system so that it meets the desired properties
- Consideration of the characteristics of a parallel system

## What this course is not about ...

- Exploiting data parallelism
- Hardware-level concurrency

# General Info

## Structure

- **Part 1:** Shared Memory (and Java)
- **Part 2:** Message Passing (and Go)
- **Part 3:** Analyses and Tool Support (and Rust)

## Literature

- First part: **G. R. Andrews. Foundations of Multi-threaded, Parallel, and Distributed Programming.** Addison Wesley, 2000 (Chapters 1 to 5, 7 and 8).
- Second and third part: Slides from the lectures, supplementary material

# General Info

## Exercises, obligs, exam

- (Almost) weekly exercise sessions (check schedule online)
- Two obligatory assignments (obligs):
  - Oblig 1: handout on 29.08.2025, submission deadline 15.09.2025
  - Oblig 2: handout on 19.09.2025, submission deadline 20.10.2025
- **Exam 11.12.2025**, check website for details

## Rooms

- **Lectures:** 10:15 on Wednesdays in Store auditorium, Kristen Nygaards hus  
(except 1<sup>st</sup> lecture which is in Seminarrom C, Ole-Johan Dahls hus)
- **Exercises:** 12:15 on Fridays in Prolog
- First exercise on Friday, 29.08.2025

# Today's Agenda

- Overview ✓
- Motivation and considerations
- Critical sections and the await-language
- A taste of concurrency in Java

**Reading material:** Chapter 1 of Andrews, additional slides about Java (see course page)

## Shared Memory Concurrency

---

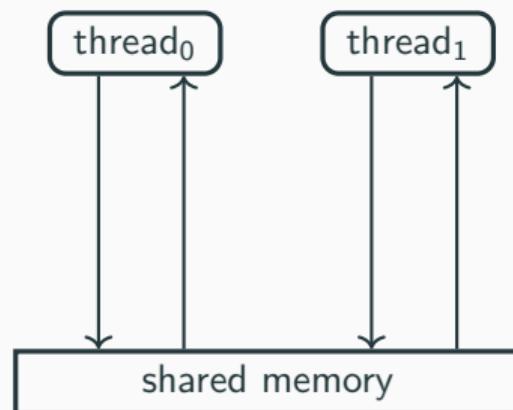
# Parallel Processes

- **Sequential program:** one thread of control, full control over the whole memory
- **Parallel/concurrent program:** several threads of control, which *need to exchange information and coordinate execution*

## Communication between processes

We will study two different ways to organize communication between processes:

- Reading from and writing to *shared variables* (Now)
- Communication with *messages* between processes (Part 2)



# Course Overview — Part 1: Shared Variables

## Content

- Problems that occur in concurrent systems with shared variables
- Patterns to solve these problems

- Atomic operations
- Interference
- Deadlock, livelock, liveness, fairness
- Locks, critical sections and (active) waiting
- Semaphores and passive waiting
- Monitors
- Java: threads and synchronization

# Why Shared Variables?

## Why shared (global) variables?

- Reflected in conventional hardware architectures, e.g., multi-core systems
- Reflected in most programming languages as a default (e.g., multi-threading)

## Notes

- Even with a single processor and one thread, you may want to use many processes, in order to get a natural partitioning, e.g., several active windows at the same time
- As all concurrency: potentially greater efficiency and/or better latency if several things happen/appear to happen “at the same time”.
- Natural interaction for tightly coupled systems

## Simple Example

Consider 3 global variables  $x$ ,  $y$ , and  $z$  and the following program:  $x := x+z$ ;  $y := y+z$ ;

- Can we use parallelism here (without changing the results)?
- If operations can be performed *independently* then performance may increase
- What are the results?

### Await

Before:  $\{x = a \& y = b \& z = c\}$

$x := x+z$ ;  $y := y+z$ ;

After:  $\{x = a+c \& y = b+c \& z = c\}$

### Pre/post-conditions

- We use brackets  $\{\dots\}$  to describe conditions before or after a statement
- These conditions are describing the state, but are not executed
- Java has `assert` that can check such conditions at runtime and the Java Modeling language (JML) to specify more complex conditions than expressions

# Parallel Operator ||

- Consider shared and non-shared program variables, assignment.
- We extend the language with a construction for *parallel composition*:

Await

```
co S1 || S2 || ... || Sn oc
```

- The execution of a parallel composition happens via the *concurrent* execution of the component processes  $S_1, \dots, S_n$ .
- *Terminates* normally if all component processes terminate normally.

## Example

Await

```
{x = a & y = b & z = c}  
x := x+z; y := y+z;  
{x = a+c & y = b+c & z = c}
```

Await

```
{x = a & y = b & z = c}  
co x := x+z || y := y+z oc  
{x = a+c & y = b+c & z = c}
```

## Simple Example in Java

Java

```
class C { public static int x = 0, y = 0, z = 5 }

...
new Thread ( () -> { x = x + z } ).start();
new Thread ( () -> { y = y + z } ).start();
```

# Interaction Between Parallel Processes

Processes can *interact* with each other in *two* different ways:

- *Cooperation* to obtain a result
- *Competition* for common resources

To organize their interactions, we use *synchronization*

## Synchronization

Synchronization *restricts* the possible interleavings of parallel processes to avoid unwanted behavior and enforce wanted behavior.

## Example

- Increasing *atomicity* and *mutual exclusion* (Mutex) to introduce *critical sections* which can *not* be executed concurrently
- *Conditional synchronization* enforces that processes must *wait* for a specific condition to be satisfied before execution can continue.

# Atomic Expressions Java Virtual Machine (JVM) bytecode

## JVM

Java is compiled down to JVM bytecode, which does not correspond 1-to-1 to machine instructions

## Java

```
x++; // x is only local variable, declared as long
```

How many machine instructions will this be? **4-6**

```
LLOAD_1      // push value from local variable #1
LCONST_1     // push value 1
LADD         // add 2 top-most values
LSTORE_1     // store value into local variable #1
```

# Atomic Expressions JVM bytecode

## JVM

Java is compiled down to JVM bytecode, which does not correspond 1-to-1 to machine instructions

## Java

```
x++; // x is only local variable, declared as long
```

How many machine instructions will this be? **4-6**

- Reference reads and writes are atomic
- Basic type reads and writes except long and double are guaranteed to be atomic
- On 64bit machines, long and double reads and writes *might* be atomic, *might* be two instructions

# Concurrent Processes: Atomic Operations

## Definition (Atomic)

An operation is **atomic** if it cannot be subdivided into smaller operations.

- We can ignore concurrency inside atomic operations as they cannot be interleaved
- A statement with at most one atomic operation, in addition to operations on **local** variables, can be considered atomic
- What is atomic depends on the language/setting:  
**fine-grained** and **coarse-grained** atomicity.
- Accessing global variables is atomic for this lecture.  
(In general, this may not be the case, e.g., for long.)
- Assignments  $x := e$  are *not* atomic

# Atomic Operations on Global Variables

Enabling atomic operations on global variables is fundamental for shared memory concurrency

- Process *communication* may be realized by variables:  
a communication channel corresponds to a variable of type vector (or similar)
- Associated with global variables is a set of *atomic operations*
- Typically: read and write, in hardware, e.g. LOAD/STORE to registers
- Atomic operations on a variable  $x$  are called  **$x$ -operations**

Our goal:

Mutual exclusion

Make *composed* statements atomic so they cannot happen simultaneously.

... but observe: the more atomic we make the program, the less parallel execution can occur!

## Example

Await

P1

P2

{ $x = 0$ } **co**  $x := x+1$  ||  $x := x-1$  **oc** {?}}

Each statement actually consists of 3 operations, e.g.,  $P_1$  is

read  $x$ ; inc; write  $x$ ;

Atomic  $x$ -operations:

- $P_1$  reads value of  $x$  (R1)
- $P_1$  writes a value into  $x$  (W1)
- $P_2$  reads value of  $x$  (R2)
- $P_2$  writes a value into  $x$  (W2)

What is the final state of our program?

## Interleaving & Possible Execution Sequences (I)

The four operations cannot be executed in any order, the *program order* gives two constraints

- R1 must happen before W1
- R2 must happen before W2

### Definition (Program Order)

- Two statements  $S_1, S_2$  are program-ordered if they are in the same thread of the program, and  $S_1$  occurs before  $S_2$ .
- Two operations  $O_1, O_2$  from the same statement are program-ordered if  $O_1$  occurs before  $O_2$  in the translation of the statement.

In the example, inc and dec (" -1 ") are process-local, so we can ignore them

## Interleaving & Possible Execution Sequences (II)

### Definition (Interleaving)

An interleaving of two sequences  $A, B$  is a sequence  $C$ , such that

- exactly all elements of  $A$  and  $B$  are elements of  $C$ , and
- the order of elements in  $A$  (resp.  $B$ ) is respected in  $C$

An interleaving may have additional constraints (for us, e.g, program order).

### Interleavings for our example:

R1	R1	R1	R2	R2	R2
W1	R2	R2	R1	R1	W2
R2	W1	W2	W1	W2	R1
W2	W2	W1	W2	W1	W1
x	0	-1	1	-1	1
					0

## Non-deterministic Behavior

- Final values for  $x$ :  $\{0, 1, -1\}$
- As (post)-condition:  $-1 \leq x \leq 1$
- Which one is chosen during an execution?

Await

```
{x = 0} co x := x+1 || x := x-1 oc {-1 ≤ x ≤ 1}
```

- **Non-determinism:** some choices for the program are decided during execution
- For us: the exact interleaving of instructions
- In practice, choices are not “random”, but depend on factors *outside* the program code:
  - Timing of the threads
  - Scheduler of the operating system
  - ...

# Execution-space Explosion

How many interleavings of statements are possible for one given input?

- Assume that we have 3 processes, each with the same number of atomic operations, and the same starting state
- Consider executions of  $P_1 \parallel P_2 \parallel P_3$

nr. of atomic operations	nr. of executions
2	90
3	1680
4	34 650
5	756 756

- Factorial growth!
- Different executions can lead to different final states.
- Even for simple systems: *impossible* to consider every possible execution in isolation

# Factorial Explosion

The number of executions grows exponentially!

For  $n$  processes with  $m$  atomic statements each:

$$\text{number of executions} = \frac{(n \cdot m)!}{m!^n}$$

- $n=m=5$  gives  $311680371562560$ , i.e.  $> 3 \cdot 10^{14}$
- It would take ten million years to check, checking one execution each second!
- ... for each choice of input
- Testing hopeless as a validation technique!

*How can we reduce the complexity?*

# The “at-most-once” Property

Fine-grained atomicity

Only the most basic operations (R/W) are atomic

- However, some non-atomic interactions appear to be atomic
- Note expressions only perform read-accesses (unlike statements)
- A *critical reference* in an expression  $e$  is a variable that is changed by another process
- An expression without critical references is evaluated as if atomic

Definition (At-most-once property)

$x := e$  satisfies the *amo*-property if either

1.  $e$  contains *no* critical reference, or
2.  $e$  contains *at most one* critical reference and  $x$  is not *referenced* by other processes

Assignments with the at-most-once property can be considered atomic!

## At-most-once Examples

x, y shared variables, and r, s local variables

Await

```
{x=y=0} co x := x+1 || y := x+1 oc {x = 1 & (y = 1 | y = 2)}  
{x=y=0} co x := y+1 || y := x+1 oc {(x,y) ∈ {(1,1),(1,2),(2,1)}}  
{x=y=0} co x := y+1 || x := y+3 || y := 1 oc {y = 1 & 1<=x<=4}  
{y=0} co r := y+1 || s := y-1 || y := 5 oc {?} 
```

Beware of unintuitive behavior:

Await

```
{ x = 0 } co r := x-x || x := 5 oc { r = 0? }  
{ x = 0 } co x := x || ... oc { ? } 
```

# A Minimal Language for Concurrency

---

# The Await Language

- Await is used to illustrate basic ideas about concurrency without boilerplate from mainstream languages
- Their implementation in mainstream languages is shown afterwards

## Features of Await

- Standard imperative constructs: sequence (;), assignment, branching, loops
- **co** ... || ... **oc** for parallel execution
- < ... > for atomic sections
- **await** for synchronization

# Syntax: The Sequential Part

We use the following syntax for non-parallel control-flow

## Declarations

```
int i = 3  
int a[n]
```

## Assignments

```
x := e  
a[i] := e  
x++  
sum +:= i
```

- **Sequential composition**  
statement; statement
- **Compound statement (block)**  
{statement}
- **Conditional**  
if ( condition ) statement
- **While-loop**  
while ( condition ) statement
- **For-loop**  
for [ i=0 to n-1 ] statement

# Parallel Statements

Await

```
co S1 || S2 || ... || Sn oc
```

- The statements  $S_i$  are executed *in parallel* with each other
- The parallel statement terminates when all  $S_i$  have terminated (“join” synchronization)

Await

```
{x = 0 & y = 0} co x := 1 || y := 1 oc; z := x+y { z = 2 }
```

# Parallel Processes

For modularity, we also allow processes

*Await*

```
process foo {
    int sum := 0;
    for [i=1 to 10]
        sum += 1;
    x := sum;
}
```

- Processes are declared globally
- All declared processes are started in the beginning and evaluated in an arbitrary order

## Example

Await

```
process bar1{  
    for [i = 1 to n]  
        write(i);  
}
```

**Starts one process**

The numbers are printed in increasing order.

Await

```
process bar2a{ write(1); }  
process bar2b{ write(2); }
```

**Starts two processes**

The numbers are printed in arbitrary order because the execution order of the processes is *non-deterministic*.

Await

```
process barn [i=1 to n]{  
    write(i);  
}
```

**Starts  $n$  processes**

The numbers are printed in arbitrary order.

# Threads in Java

Java

```
class Printer implements Runnable {  
    private String text;  
    public Printer(String text) { this.text = text; }  
    public void run() { System.out.println(text); }  
    public static void main(String args[]) {  
        Thread t1 = new Thread(new Printer("Hello"));  
        Thread t2 = new Thread(new Printer("Concurrency"));  
        t1.start(); t2.start();  
    }  
}
```

- The Thread class encapsulates a system thread
- The Runnable interface is used to define thread behavior

# Synchronization in Java

- Java does not have atomic blocks in the same form
- Instead: synchronized methods and blocks

Java

```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() {c++;}  
    public synchronized void decrement() {c--;}  
    public synchronized int value() {return c;}  
}
```

- Synchronized methods are atomic *per object*
- Only one thread can execute a synchronized method at any time and no other thread can execute any other synchronized method on the same object

## Java and Await

Await

```
co s1 || s2 oc; s3
```

Java

```
Thread t1 = new Thread(() -> {s1});  
Thread t2 = new Thread(() -> {s2});  
Thread t3 = new Thread(() -> {s3});  
t1.start(); t2.start();  
t1.join(); t2.join();  
t3.start();
```

# Java and Await

Await

```
co <s1> || <s2> oc
```

Java

```
public class C() {  
    public static synchronized void m1(){ s1 }  
    public static synchronized void m2(){ s2 }  
    ...  
  
    new Thread(() -> {C.m1();}).start();  
    new Thread(() -> {C.m2();}).start();  
}
```

## Semantic Concepts (“Interleaving Semantics”)

- A *state* in a parallel program consists of the values of the variables at a given moment in the execution.
- Each process executes independently of the others by *modifying* global variables using atomic operations.
- Do we really need to consider all interleavings to reason about possible states?
  - How to exclude some interleavings?
  - How to make reasoning modular and compositional?

Next, a first helping concept: interference

## Read- and Write-variables

- $\mathcal{V}$  : statement  $\cup$  expression  $\rightarrow \mathcal{P}(\text{variable})$ : set of global variables in a statement or expression
- $\mathcal{W}$ : statement  $\cup$  expression  $\rightarrow \mathcal{P}(\text{variable})$ : set of global *write*-variables

$$\mathcal{V}(x := e) = \mathcal{V}(e) \cup \{x\}$$

$$\mathcal{V}(S_1; S_2) = \mathcal{V}(S_1) \cup \mathcal{V}(S_2)$$

$$\mathcal{V}(\text{if } (b) \text{ then } S) = \mathcal{V}(b) \cup \mathcal{V}(S)$$

$$\mathcal{V}(\text{while}(b)S) = \mathcal{V}(b) \cup \mathcal{V}(S)$$

Remaining cases analogous

$\mathcal{W}$  analogously, except the only difference for read-only expressions.

$$\mathcal{W}(\text{expression}) = \emptyset$$

### Example

$$\mathcal{W}(x := e) = \{x\}$$

# Disjoint Processes

## Interference freedom

Processes without common global variables are *interference-free*

$$\mathcal{V}(S_1) \cap \mathcal{V}(S_2) = \emptyset$$

- Statements obviously cannot perform any action that influence each other
  - As all interleavings are the same, one can just run  $S_1; S_2$  for analysis
  - Sequence  $S_1; S_2$  is of course less performant
- 
- If variables accessed by both processes are *read-only* variables, the same holds
  - Is the following *interference criterion* sufficient?

$$\mathcal{W}(S_1) \cap \mathcal{W}(S_2) = \emptyset$$

- Write-variables are important for *race conditions*, *critical references*/amo-property, ...
- If only read-only variables are accessed, no races or critical references exist

## Critical Sections and Invariants

---

# Properties

Read-only variables are a very coarse way to think,  
how to express more specific properties?

- A *property* is a predicate over a program, resp. its execution and reachable states
- A program has a property, if the property is true for all possible executions of the program

## Classification (I)

- **Safety** property: program will never reach an undesirable state
- **Liveness** property: program will eventually reach a desirable state

## Classification (II)

- *Termination*: all executions are finite.
- *Partial correctness*: If the program terminates, it is in a desired final state (safety property).
- *Total correctness*: The program terminates and is partially correct.

# How to Check Properties of Programs?

- *Testing* or *debugging* increases confidence in a program, but gives no guarantee of correctness.
- *Operational reasoning* considers *all* executions of a program explicitly
- *Formal analyses* reason about the properties of a program without considering the executions one by one.

Dijkstra's dictum:

A test can only show errors, but never prove that a program is correct!

## Properties: Invariants (I)

### Definition (Invariant)

An *invariant* is a property of program states, that holds for all reachable states of a program.

- *Invariant* (adj): constant, unchanging
- Prototypical safety property
- Appropriate for non-terminating systems (does not require a final state)
- All reachable states often too strong

### Kinds of invariants

- Strong invariant: Holds for all reachable states
- Weak invariant: Holds for all states where an atomic block starts or ends
- Loop invariant: Holds at the start and end of a loop body
- Global invariant: Reasons about state of many processes
- Local invariant: Reasons about state of one process

## Properties: Invariants (II)

- How to show that a program has a weak invariant?
- Without exploring all executions?

### Induction for invariants

One can show that a program has a weak invariant by

1. Showing that the invariant property holds initially,
2. and that each atomic statement maintains the property

# Critical Sections

To enforce atomicity, we have a special construct in the language :  $\langle S \rangle$  performs  $S$  atomically

## Use of critical sections

- When the processes interfere: *synchronization* to restrict the possible interleavings
- Synchronization gives coarser grained atomic operations (“atomic blocks”)
- Combines operations into an *atomic block* where the process shall not be interrupted

## Characteristics of atomic operations

- Internal states are *not visible* to other processes.
- Variables *cannot* be changed underway by other processes.
- $S$ : executed like a *transaction*

## Await

```
int x:=0; co <x:=x+1> || <x:=x-1> oc {x=0}
```

# Conditional Critical Sections

## Await statement

The **<await (B) S>** statement executes the statement *S* once the boolean condition *B* holds.

- Boolean condition *B*: *await condition*, evaluated atomically
- Body *S*: *critical section* executed atomically

The following delays the decrement until  $y > 0$  holds – or does not terminate if it never holds

### *Await*

```
<await (y > 0) y := y-1> { y >= 0 }
```

- Important that *B* has no side-effects!

## Typical Pattern for Critical Sections

- One wants to avoid using atomic blocks as much as possible
- Use them in certain places to enable correct, interleaved executions

### Await

```
int counter = 1; // global variable

// start CS
< await (counter > 0) counter := counter -1; >
critical statements;
// end CS
counter := counter +1
```

- “Critical statements” *not* enclosed in atomic block
- Invariant:  $0 \leq \text{counter} \leq 1$  (= counter acts as *binary lock*)
- Next lectures: patterns for correctness while minimizing atomic blocks

## Example: Synchronization of Strongly Coupled Producer-Consumer System

Await

```
int buf, p := 0; c := 0;
```

Await

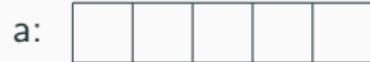
```
process Producer {  
    int a[N]; ...  
    while(p < N){  
        <await (p = c) >;  
        buf := a[p];  
        p := p+1;  
    }  
}
```

Await

```
process Consumer {  
    int b[N]; ...  
    while (c < N) {  
        <await (p>c) >;  
        b[c] := buf;  
        c := c+1;  
    }  
}
```

- buf as only shared variable, acting as a one element buffer
- Tasks of synchronization:
  - Coordinating the “speed” of the two processes
  - Avoiding to read data which is not yet produced

## Example (Continued)



buf: A single empty square box.  
p: A single empty square box.  
c: A single empty square box.  
N: A single empty square box.



- A strong invariant holds in *all states* in *all* executions of the program.
- *Global invariant*:  $c \leq p \leq c+1$
- *Local invariant (Producer)*:  $0 \leq p \leq N$

# Wrap-Up

## Summary

- Shared memory
- Synchronization
- Atomic operations,
- Interleavings
- await-language and critical sections
- A taste of concurrency in Java