

Semaphores

Silvia Lizeth Tapia Tarifa

September 10, 2025

University of Oslo

Overview

Last lecture: Locks and Barriers

- Complex techniques
- No clear separation between variables for synchronization and variables for computation
- Busy waiting

This lecture: Semaphores

- Synchronization tool
- Used easily for mutual exclusion and condition synchronization
- A way to implement signaling and scheduling
- Implementable in many ways on hardware (CMPXCHG)
- Available in programming language libraries and OS

Outline


- Semaphores: Syntax and semantics
- Synchronization examples:
 - Mutual exclusion (critical sections)
 - Barriers (signaling events)
 - Producers and consumers (split binary semaphores)
 - Bounded buffer: resource counting
 - Dining philosophers: mutual exclusion – deadlock
 - Readers and writers:
 - condition synchronization
 - passing the baton

Semaphores

Semaphores

Origins of Term

- Introduced by Dijkstra in 1968
- Inspired by railroad traffic synchronization
- Railroad semaphore indicates whether the track ahead is clear or occupied by another train



Properties

- Semaphores in concurrent programs: work similarly
- Used to implement
 - *mutex* and
 - *condition synchronization*
- Included in most standard libraries for concurrent programming
- Also *system calls* in, e.g., Linux kernel, Windows etc.

Concept

Concept of a Semaphore

- *Semaphore*: special kind of shared program variable (with built-in sync. power)
- value of a semaphore: a *non-negative* integer
- can *only* be manipulated by two *atomic* operations:

The Semaphore Operations: *P* and *V*

- **P**: (Passeren) Wait for signal – want to *pass*
Wait until value is greater than zero, and *decrease* value by one
- **V**: (Vrijgeven) Signal an event – *release*
Increase the value by one

- Today, libraries and sys-calls prefer other names: up/down, wait/signal, acquire/release
- Different flavors of semaphores: binary vs. counting
- Most common: mutex as a synonym for binary semaphores

Syntax and Semantics

Declaration

- sem s; default initial value is zero
- sem s := 1;
- sem s[4] := ([4] 1);

Operations and Semantics

P-operation P(s)

$\langle \text{await } (s > 0) s := s - 1 \rangle$

V-operation V(s)

$\langle s := s + 1 \rangle$

Processes waiting on a semaphore are woken up by the op. system.

Remarks on Semaphores

Remark 1

Important: No *direct* access to the value of a semaphore.

For example, a test like if $(s = 1)$ then ... else is *forbidden*!

Kinds of semaphores

General semaphore: Possible values: *all non-negative integers*

Binary semaphore: Possible values: 0 and 1

Fairness

- As for await-statements.
- In most languages: *FIFO* (“waiting queue”): processes delayed while executing P-operations are *awoken* in the *order* they were delayed

Example: Mutual Exclusion (critical section)

Mutex implemented by a *binary semaphore*

Await

```
sem mutex := 1;
process CS[ i = 1 to n] {
    while (true) {
        P(mutex);
        # critical section
        V(mutex);
        # noncritical section
    }
}
```

- The semaphore is *initially 1*
- Always P before V → (used as) binary semaphore

Example: Barrier Synchronization


Semaphores may be used for *signaling events*

Await

```
sem arrive1 = 0, arrive2 = 0;
process Worker1 {
    ...
    V(arrive1); # reach barrier
    P(arrive2); # wait for other
    ...
}
process Worker2 {
    ...
    V(arrive2); # reach barrier
    P(arrive1); # wait for other
    ...
}
```

Example: Barrier Synchronization

- *Signalling semaphores*: usually *initialized to 0* and
- *Signal* with a V and then *wait* with a P



Split Binary Semaphores

Split binary semaphore

A set of semaphores, whose *sum* ≤ 1

Mutex by split binary semaphores

- Initialization: *one* of the semaphores = 1, all others = 0
- Discipline: all processes call *P* on a semaphore, *before* calling *V* on (*another*) semaphore
⇒ Code between the *P* and the *V*
 - All semaphores = 0
 - Code executed *in mutex*

Example: Producer/Consumer with Split Binary Semaphores

Await

```
T buf; # one element buffer, some type T  
sem empty := 1;  
sem full := 0;
```

Await

```
process Producer {  
    while (true) {  
        P(empty);  
        buff := data;  
        V(full);  
    }  
}
```

Await

```
process Consumer {  
    while (true) {  
        P(full);  
        data_c := buff;  
        V(empty);  
    }  
}
```

- empty and full are both *binary semaphores*, together they form a split binary semaphore.
- Solution works with *several* producers/consumers


Example: Producer/Consumer with Split Binary Semaphores

Await

```
process Producer {  
    while (true) {  
        P(empty);  
        buff := data;  
        V(full);  
    }  
}
```


Await

```
process Consumer {  
    while (true) {  
        P(full);  
        data_c := buff;  
        V(empty);  
    }  
}
```



Producer/Consumer: Increasing Buffer Capacity

- Previously: tight coupling, the producer must wait for the consumer to empty the buffer before it can produce a new entry.
- Easy *generalization*: buffer of size n .
- Loose coupling/asynchronous communication \Rightarrow “buffering”
 - *Ring-buffer*, typically represented
 - by an array
 - + two integers **rear** and **front**.
 - Semaphores to *keep track* of the number of free/used slots \Rightarrow general semaphore



Producer/Consumer: Increased Buffer Capacity

Await

```
T buf[n]                      # array , elements of type T
int front := 0, rear := 0; # ``pointers''
sem empty := n;             # number of empty slots
sem full := 0;              # number of filled slots
```

Await

```
process Producer {
    while (true) {
        P(empty);
        buff[rear] := data;
        rear := (rear + 1);
        V(full);
    }
}
```

Await

```
process Consumer {
    while (true) {
        P(full);
        result := buff[front];
        front := (front + 1);
        V(empty);
    }
}
```

Producer/Consumer: Increased Buffer Capacity


Await

```
process Producer {
    while (true) {
        P(empty);
        buff[rear] := data;
        rear := (rear + 1);
        V(full);
    }
}
```

Await

```
process Consumer {
    while (true) {
        P(full);
        result := buff[front];
        front := (front + 1);
        V(empty);
    }
}
```

- Important: there are no critical sections!
- How to enable several producers and consumers?



Increasing the Number of Processes

How to enable several producers and consumers?

New synchronization problems

- *Avoid* that two producers *deposit* to buf [rear] before rear is updated
- *Avoid* that two consumers *fetch* from buf [front] before front is updated.

Solution

Add 2 extra binary semaphores for protection:

- mutexDeposit to deny two producers to deposit to the buffer at the same time.
- mutexFetch to deny two consumers to fetch from the buffer at the same time.

Example: Producer/Consumer with Several Processes

Await

```
T buf[n]                      # array, elements of type T
int front := 0, rear := 0; # ``pointers''
sem empty := n;             # number of empty slots
sem full := 0;              # number of filled slots
sem mutexDeposit; mutexFetch := 1; # protect the data stuct.
```


Await

```
process Producer {
    while (true) {
        P(empty);
        P(mutexDeposit);
        buff[rear] := data;
        rear := (rear + 1);
        V(mutexDeposit);
        V(full);
    }
}
```

Await

```
process Consumer {
    while (true) {
        P(full);
        P(mutexFetch);
        result := buff[front];
        front := (front + 1);
        V(mutexFetch);
        V(empty);
    }
}
```

Problem: Dining Philosophers



source:wikipedia.org

- Famous sync. problem (Dijkstra)
- Five philosophers around a circular table.
- One fork placed between each pair of philosophers
- Each philosopher alternates between thinking and eating
- A philosopher needs two forks to eat (and none for thinking)

Dining Philosophers: Sketch

Await

```
process Philosopher [i = 0 to 4] {
    while true {
        # think
        acquire forks;
        # eat
        release forks;
    }
}
```

Task:


Program the actions `acquire forks` and `release forks`

Dining philosophers: 1st attempt

- Forks as *semaphores*
- Philosophers: pick up left fork first

Await

```
sem fork[5] := ([5] 1);
process Philosopher [ i = 0 to 4 ] {
    while true {
        # think
        P(fork[i]);
        P(fork[(i+1)%5]);
        # eat
        V(fork[i]);
        V(fork[(i+1)%5]);
    }
}
```



Dining philosophers: 2nd attempt

Breaking the symmetry

To avoid *deadlock*, let 1 philosopher (say 4) grab the *right* fork first

Await

```
process Philosopher [ i = 0 to 3] {
    while true {
        think;
        P(fork[i]);
        P(fork[(i+1)%5]);
        eat;
        V(fork[i]);
        V(fork[(i+1)%5]);
    }
}
```

Await

```
process Philosopher4 {
    while true {
        think;
        P(fork[4]); #!
        P(fork[0]); #!
        eat;
        V(fork[4]);
        V(fork[0]);
    }
}
```

Dining philosophers

- Important illustration of problems with concurrency:
 - Deadlocks,
 - Other aspects: liveness, fairness, etc.
- Resource access
- Connection to mutex/critical sections

Invariants and Condition Synchronization

Readers/Writers: Overview

- Classic synchronization problem
- *Reader* and *writer* processes, share access to a database/shared data structure
 - Readers only read from the database
 - Writers update (and read from) the database
- As soon as one writer is included, read and write accesses may cause interference
- Readers and writers have **asymmetric requirements:**
 - Every *writer* needs *mutually exclusive* access
 - When no writers have access, *many readers* may access the database

Readers/Writers: Approaches

- Dining philosophers: Pair of processes compete for access to “forks”
- Readers/writers: Different *classes* of processes compete for access to the database
 - Readers *compete* with writers
 - Writers *compete* both with readers and other writers
- **General synchronization problem:**
 - Readers: must wait until no writers are active in DB
 - Writers: must wait until no readers or writers are active in DB
- Here: two different approaches
 1. **Mutex:** easy to implement, but “*unfair*”
 2. **Condition synchronization:**
 - Using a *split binary semaphore*
 - Easy to adapt to different scheduling strategies

Readers/Writers with Mutex (1)

Await

```
sem rw := 1;
```

Await

```
process Reader [ i=1 to M] {
    while (true) {
        P(rw);
        # read
        V(rw);
    }
}
```

Await

```
process Writer [ i=1 to N] {
    while (true) {
        P(rw);
        # write
        V(rw);
    }
}
```

We want *more than one reader simultaneously*.

Readers/Writers with Mutex (2)

Await

```
int nr := 0;    # number of active readers
sem rw := 1      # lock for reader/writer mutex
```

Await

```
process Reader [ i=1 to M] {
    while (true) {
        < nr := nr + 1;
        if (nr=1) P(rw) >;
        # read
        < nr := nr - 1;
        if (nr=0) V(rw) >;
    }
}
```

Await

```
process Writer [ i=1 to N] {
    while (true) {
        P(rw);
        # write
        V(rw);
    }
}
```

How do semaphores work *inside* atomic sections?

Readers/Writers with Mutex (3)

Await

```
int      nr = 0; # number of      active readers
sem      rw = 1; # lock for reader/writer exclusion
sem mutexR = 1; # mutex for readers

process Reader [ i=1 to M] {
    while (true) {
        P(mutexR)
        nr := nr + 1;
        if (nr=1)  P(rw);
        V(mutexR)
        # read
        P(mutexR)
        nr := nr - 1;
        if (nr=0)  V(rw);
        V(mutexR)
    }
}
```

Readers/Writers with Condition Synchronization: Overview

Reader's preference

- With a constant stream of readers, the writer will never run
- Even under strong fairness
- Previous *mutex* solution solved *two* separate synchronization problems
 - ***rw*** : *Readers and writers* for access to the *database*
 - ***mutexR***: *Reader vs. reader* for access to the *counter*
- Now: a solution based on **condition synchronization**

Invariant

Reasonable invariant for the critical sections

1. When a *writer* accesses the DB, *no one else* can
2. When *no writers* access the DB, *one or more readers* may get access

Introducing state for the invariant

Introduce two counters:

- **nr**: number of active readers
- **nw**: number of active writers

Invariant

$RW: \quad (nr = 0 \text{ or } nw = 0) \text{ and } nw \leq 1$
(same as:) $RW': \quad nw=0 \text{ or } (nw = 1 \text{ and } nr = 0)$

Code for counting Readers and Writers

Await

```
< nr := nr + 1; >  
# read  
< nr := nr - 1; >
```

Await

```
< nw := nw + 1; >  
# write  
< nw := nw - 1; >
```

- Add synchronization code to maintain the invariant
- Decreasing counters is not dangerous
- Before increasing, we need to check some conditions for synchronization
 - before increasing nr: nw = 0
 - before increasing nw: nr = 0 and nw = 0

Condition Synchronization: Without Semaphores

Await

```
int nr := 0;    # number of active readers
int nw := 0;    # number of active writers
# Invariant RW: (nr = 0 or nw = 0) and nw <= 1
```

Await

```
process Reader [ i=1 to M]{
  while (true) {
    < await (nw=0)
      nr := nr+1>;
    # read
    < nr := nr - 1>
  }
}
```

Await

```
process Writer [ i=1 to N]{
  while (true) {
    < await (nr = 0 and nw = 0)
      nw := nw+1>;
    # write
    < nw := nw - 1>
  }
}
```

Condition Synchronization: Converting to Split Binary Semaphores

Convert awaits with different guards $B_1, B_2 \dots$ to Split Binary Semaphores

- *Entry:* semaphore e , initialized to 1
- For each guard B_i :
 1. associate one *counter* and
 2. one *delay-semaphore*

Both initialized to 0

- Semaphore *delays* the processes waiting for B_i
- Counters counts the number of processes *waiting* for B_i

For *readers/writers* problem we need 3 semaphores and 2 counters:

Await _____

```
sem e = 1;
sem r = 0; int dr = 0; # condition reader: nw == 0
sem w = 0; int dw = 0; # condition writer: nr == 0 and nw == 0
```

Condition Synchronization: Converting to Split Binary Semaphores (2)

- e , r and w form a *split binary semaphore*.
- All execution paths *start* with a *P-operation* and *end* with a *V-operation* → Mutex

Signaling

We need a signal mechanism *SIGNAL* to pick which semaphore to signal.

- SIGNAL: make sure the invariant holds
- B_i holds when a process enters *CS* because either:
 - the process checks itself,
 - or the process is only *signalled* if B_i holds
- **Another pitfall:**

Avoid *deadlock* by checking the counters before the delay semaphores are signalled.

 - r is not signalled ($V(r)$) *unless* there is a delayed reader
 - w is not signalled ($V(w)$) *unless* there is a delayed writer

Condition Synchronization: Reader

Await

```
int nr := 0, nw = 0;      # counter variables (as before)
sem e := 1;                # entry semaphore
int dr := 0; sem r := 0; # delay counter + sem for reader
int dw := 0; sem w := 0; # delay counter + sem for writer
# invariant RW: (nr = 0 or nw = 0 ) and nw <= 1
process Reader [i=1 to M]{ # entry condition: nw = 0
    while (true) {
        P(e);
        if (nw > 0) { dr := dr + 1;    # < await (nw=0)
                        V(e);          #     nr:=nr+1 >
                        P(r)};
        nr:=nr+1; SIGNAL;
        # read
        P(e); nr:=nr-1; SIGNAL;      # < nr:=nr-1 >
    }
}
```

With Condition Synchronization: Writer

Await

```
process Writer [i=1 to N]{ # entry condition: nw = 0 and nr = 0
    while (true) {
        P(e);                                # < await (nr=0 and nw=0)
        if (nr > 0 or nw > 0) {      #     nw:=nw+1 >
            dw := dw + 1;
            V(e);
            P(w) };
        nw:=nw+1; SIGNAL;
        # write
        P(e); nw:=nw-1; SIGNAL      # < nw:=nw-1>
    }
}
```

With Condition Synchronization: Signalling

Await

```
if (nw = 0 and dr > 0) {  
    dr := dr -1; V(r);           # awake reader  
}  
elseif (nr = 0 and nw = 0 and dw > 0) {  
    dw := dw -1; V(w);         # awake writer  
}  
else V(e);                      # release entry lock
```

- This passes the control (the “baton”) to an appropriate next process
- SIGNAL has no P operation, each path has exactly one V operation.
- Using the conditions to see who goes next.
- Called “passing the baton” technique (as in relay competition).
- Conditions for awakening must be disjoint

Example: 1 Reader, 1 Writer, Reader starts

| | | | | | | | | | | | | | | | | | |
|----|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| nr | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| nw | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 0 | 0 |
| e | 1 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 0 |
| dw | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| w | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Await

```
if (nw = 0 and dr > 0) {
    dr := dr -1; V(r);                      # awake reader
}
elseif (nr = 0 and nw = 0 and dw > 0) {
    dw := dw -1; V(w);                      # awake writer
}
else V(e);                                # release entry lock
```

Await

```
process Reader [i=1 to M]{  # entry condition: nw = 0
    while (true) {
        P(e);
        if (nw > 0) { dr := dr + 1;  # < await (nw=0)
```

Semaphores in Java

Basic Methods of Semaphores in Java

- `Semaphore(int n)`
 - constructor for semaphores
 - initializes semaphore value with integer *n set of permits*
- `acquire()`
 - corresponds to the P operation
 - tries to decrease the number of permits by 1
 - blocks, if that is not possible and waits, until semaphore gives permit
- `release()`
 - corresponds to the V operation
 - increases the number of permits by 1

Dining Philosophers: Naïve Solution in Java (I)

Philosophers in Java

- Philosopher has references to two binary Semaphores (leftFork and rightFork),
- and the functions eat(), sleep() and run()

Java

```
Semaphore[] forks = new Semaphore[numberOfPhilosophers];
for (int i=0; i < forks.length; i++)
    forks[i] = new Semaphore(1);

philosophers = new Philosopher[numberOfPhilosophers];
for (int i=0; i < philosophers.length; i++)
    philosophers[i] =
        new Philosopher(i, forks[i], forks[(i+1) % forks.length]);
```

Dining Philosophers: Naïve Solution in Java (II)

Java

```
while(true) {
    think();                                // think
    if(i == 0) {
        rightFork.acquire();                // acquire forks
        leftFork.acquire();
    } else {
        leftFork.acquire();                // acquire forks
        rightFork.acquire();
    }
    eat();                                    // eat
    leftFork.release();                     // release forks
    rightFork.release();
}
```

The Condition Interface

- A **condition** allows to transfer the ownership of the lock without lock/unlock
- Each condition is, thus, bound to a lock

The Condition interface includes the following methods:

- `cond.await()`
 - The lock associated with the Condition is atomically released (unlock) and the thread becomes disabled
 - After cond is signalled, the thread continues with its instructions.
- `cond.signal()`
 - Wakes up one thread that is waiting on this Condition
- Note: threads interacting with cond still need to acquire and release its lock!

```
Lock mutex = new ReentrantLock();
Condition condition = mutex.newCondition();

public void waitingThread() throws InterruptedException {
    mutex.lock();          // thread acquires the lock
    try {
        while(/*not finished*/) {
            condition.await();      // Release the lock and wait for signal
            /* thread does something (1) */
        }
    } finally {
        mutex.unlock();    // thread releases the lock
    }
}
```

Java

```
Lock mutex = new ReentrantLock();
Condition condition = mutex.newCondition();

public void signallingThread() throws InterruptedException {
    mutex.lock();                      // thread acquires the lock ;
    try {
        /* thread does something (2) */
        condition.signal();           // Signal (wake up) one waiting thread
    } finally {
        mutex.unlock();              // thread releases the lock
    }
}
```

Producer Consumer with Locks and Conditions

Example can be found at: [Example link](#)

Conclusion

Condition synchronization

- One semaphore to protect shared variables (the counters)
- For each condition: a semaphore + a “delay” counter
- On entry: increase delay counter if your condition is not true
- Wait on your condition semaphore
- Decide who is next (SIGNAL) using
 - the conditions, and
 - the delay counters to see who is waiting to enter
- SIGNAL whenever someone should get a chance to enter.