

Message Passing and Channels

Andrea Pferscher

September 24, 2025

University of Oslo

Message Passing

Structure

- Part 1: Shared Memory (and Await)
- **Part 2:** Message Passing (and Go)
- Part 3: Analyses and Tool Support (and Rust)

Content of next part:

- Synchronous and asynchronous message passing
- Channels, actors, go-routines, asynchronous programming


Outline today

- *Asynchronous message passing*: channels, messages, primitives
- Example: filters and sorting networks
- Comparison of message passing and monitors
- *Basics synchronous message passing*

Concurrent Programming: Shared State vs. Messages

Concurrent programming

- *Concurrent program*: two or more processes that work together to perform a task.
- The processes work together by communicating with each other using:
 - *Shared variables*: One process writes into a variable that is read by another.
 - *Message passing*: One process sends a message that is received by another



Program Synchronization (Recap)

Two kinds of synchronization approaches (regardless of the form of communication)

- Mutual exclusion (mutex)
 - A program mechanism that prevents processes from accessing a shared resource at the same time.
 - Only one process or thread owns the mutex at a time.
- Condition synchronization
 - Delay a process until a given condition is true.
- To prevent race condition: when concurrent processes access and change a shared resource.
- Used for critical section.

Recap

- So far: shared variable programming
- **Now:** Distributed programming

Distributed Systems

Shared Memory vs. Distributed Memory


System architectures with shared memory:

- Many processors access the same physical memory
- Examples: laptops, fileservers with many processors on one motherboard

Distributed memory architectures:



- Each processor has private memory, communication over connections in a “network”
- Examples:
 - Multicomputer: asynchronous multi-processor with distributed memory
 - Workstation clusters: PC's in a local network, NFS (Network File System)
 - Grid system: machines on the Internet, resource sharing
 - Cloud computing: cloud storage service
 - NUMA-architectures
 - Cluster computing ...

Shared Memory Concurrency in the Real World



- Shared memory architecture is a simplification
- Out-of-order executions:
 - Due to complex memory hierarchies, caches, buffers,...
 - Due to weak memory, micro-ops, compiler optimizations,...

SMP (Symmetric Multiprocessing), Multi-Core Architecture, and NUMA



Concurrent vs. Distributed Programming

Shared-Memory Systems

- Processors share one memory
- Processors communicate via reading and writing of shared variables

Concurrent programming provides primitives to synchronize over memory

Distributed Systems

- Memory is distributed: processes cannot share variables/memory locations
- Processes communicate by sending and receiving *messages* via e.g., shared *channels*,
- or (in future lectures): communication via *RPC* and *rendezvous*

Distributed programming provides primitives to communicate


- Some concepts from distributed systems are also useful abstractions for shared memory
- Abstractions can be decoded to different primitives, e.g., channels as shared-memory
- Also: mixed shared-distributed systems

Synchronous and Asynchronous Message Passing


Message Passing

- *Message passing* refers to the sending of a message to a process.
- This message can be used to invoke a process
- Two types of message passing:
 - *Synchronous* message passing
 - *Asynchronous* message passing

Synchronous Message Passing: High Level Concept



Asynchronous Message Passing: High Level Concept



Synchronous vs. Asynchronous Message Passing: Trade Off

Synchronous message passing

- No memory buffer is required
- Concurrency is reduced
- Programs are more prone to deadlock

Asynchronous message passing

- Memory buffer is required (memory is cheap)
- Have more concurrency
- Programs are less prone to deadlock

We will comeback to this comparison later in the lecture.

Channels

Asynchronous Message Passing: Channel Abstraction

Channel

Abstraction, e.g., of a physical communication network, for one-way communication between two entities (similar to producer-consumer). For us:

- Unbounded FIFO (queue) of waiting messages
- Preserves message order
- Atomic access
- Error-free
- Typed

Numerous variants exists in different language: untyped, lossy, unnamed, bounded . . .
We will look at more complex types later

Asynchronous Message Passing: Primitives

Channel declaration

Await

```
chan c(type1 id1, ..., typeN idN);
```

Messages are n -tuples of respective types.

Communication primitives

- `send c(expr1, ..., exprN);`

Non-blocking, i.e. asynchronous: message is sent and process continues its execution

- `receive c(v1, ..., vN);`

Blocking: receiver process waits until message is sent on the channel


Message stored in variables $v1, \dots, vN$.

- `empty(c);`

True if channel is empty

Example: Message Passing

$(x,y) = (1,2)$




Await

```
chan foo(int);  
  
process A {  
    send foo(1);  
    send foo(2);}  
  
process B {  
    int x; int y;  
    receive foo(x);  
    receive foo(y);}
```

Example: Shared Channel

$(x,y) = (1,2)$ or $(2,1)$



Await

```
chan foo(int);
process A1 {
    send foo(1); }

process A2 {
    send foo(2); }

process B {
    int x; int y;
    receive foo(x);
    receive foo(y); }
```

Asynchronous Message Passing and Semaphores

A channel acts as a semaphore, where sending and receiving have the same asymmetry as **V** (increase the value of the semaphore by one) and **P** (wait until value of the semaphore is greater than zero, and then decrease the value by one).

Comparison with general semaphores

channel	\approx	semaphore
send	\approx	V
receive	\approx	P


The value of the message plays no role for the semaphore-interpretation.

Filters: One-Way Interaction

Filters **F**

A filter **F** is a process which:

- Receives messages on input channels,
- Sends messages on output channels, such that
- the output is a function of the input (and the initial state).



- Some computations are naturally seen as a composition of filters:
- stream processing, feedback loops and *dataflow programming*

Example: A Single Filter Process

Task: Sort a list of n numbers into ascending order.

Filter

Process **Sort** with input channel `input` and output channel `output`.

Example implementation: get n over `input`, then read n times from `input` and send the sorted list at once over `output`.

Sort predicate

- n : number of values sent to output.

$sent[i]$: i 'th value sent to output, $received[j]$: j 'th value received in input,

$$\forall i : 1 \leq i < n. (sent[i] \leq sent[i + 1]) \wedge$$

$$\forall i : 1 \leq i \leq n. \exists j : 1 \leq j \leq n. sent[i] = received[j] \wedge$$

$$\forall i : 1 \leq i \leq n. \exists j : 1 \leq j \leq n. received[i] = sent[j]$$

Filter for Merging of Streams

Task: Merge two sorted input streams into one sorted stream.

Process Merge with input channels in_1 and in_2 and output channel out :

$$\text{in}_1 : \langle 1 \ 4 \ 9 \dots \rangle \quad \text{in}_2 : \langle 2 \ 5 \ 8 \dots \rangle \quad \text{out} : \langle 1 \ 2 \ 4 \ 5 \ 8 \ 9 \dots \rangle$$

Special value **EOS** marks the end of an input, but result should be output online.

Merge predicate

n : number of values sent to out so far, $\text{sent}[n]$: i 'th value sent to out so far.

The following shall hold when **Merge** terminates:

$$\text{empty}(\text{in}_1) \wedge \text{empty}(\text{in}_2) \wedge \text{sent}[n + 1] = \text{EOS}$$

$$\wedge \quad \forall i : 1 \leq i < n \cdot \text{sent}[i] \leq \text{sent}[i + 1]$$

\wedge values sent to out are an *interleave* of values from in_1 and in_2

Await

```
chan in1(int), in2(int), out(int);


process Merge {
    int v1, v2;
    receive in1(v1);           # read the first two
    receive in2(v2);           # input values

    while (v1 != EOS and v2 != EOS) {
        if (v1 <= v2) { send out(v1); receive in1(v1); }
        else           { send out(v2); receive in2(v2); }
    }

    while (v1 != EOS) { send out(v1); receive in1(v1); }
    while (v2 != EOS) { send out(v2); receive in2(v2); }
    send out(EOS);
}
```

Sorting Network

To scale, we can now build a network that sorts n numbers, using a **collection** of **Merge** processes with tables of shared input and output channels.



Call-Backs to a Channel

- How to communicate a result back via channels?
- For example: Assume a process that adds two numbers it receives via a channel and then returns the result to the same channel.

Bi-directional channel

Await _____

```
chan c(int);  
process P { int a, b; receive c(a); receive c(b); send c(a+b); }
```

Requires same channel type for input and result.

Call-Backs to a Channel

- How to communicate a result back via channels?
- For example: Assume a process that adds two numbers it receives via a channel and then returns the result to a channel.

Answer channel per sender

Await _____

```
chan c(int), chan d[n](int);
process P { int a, b; int id;
    receive c(a); receive c(b); receive c(id); send d[id](a+b); }
```

Requires pre-sharing of channels, rather static.

Call-Backs to a Channel

- How to communicate a result back via channels?
- For example: Assume a process that adds two numbers it receives via a channel and then returns the result to the a channel.

Call-back channel

Await

```
chan c(...);
process P {
    int a, b;
    chan res(int);
    receive c(a); receive c(b); receive c(res);
    send res(a+b);
}
```

Requires (a) sending channels over channels and (b) more complex type for c.

Message Passing

Client-Server Applications using Messages

Roles

- Server process: repeatedly handling requests from clients
- Client processes: send requests to server, retrieve results later

Await

```
chan request(int, T1); # client ID, arguments of the operation  
chan reply[n](T2); # result of the operation
```

Await

```
process Client[i = 1 to n]{  
    ...  
    send request(i, args);  
    receive reply[i](var);  
    ...  
}
```

Await

```
process Server{  
    while(true){ int id; ...  
        receive request(id, args);  
        ... # code of the operation  
        send reply[id](result);  
    } }
```

Monitor Implementation using Message Passing

Monitors are very useful in a shared-memory setting, can we implement it in a channel-based concurrency model?

Classical monitor

- Controlled access to shared resource
- Global variables safeguard the resource state
- Access to a resource via procedures
- Procedures are executed under mutual exclusion
- Condition variables for synchronization

Active Monitors

- One server process that actively runs a loop listens on a channel for requests
- Procedure calls correspond to values send over request channel
- Resource and variables are local to the server process

Allocator for Multi-Unit Resources

Task

Multi-unit resource: a resource consisting of multiple units, which can be allocated separately, e.g., memory blocks, file blocks, etc.

- Client can request resources, use them, and return/free them
 - All the access to resources is managed for safety by the allocator
 - Unit usage itself is not managed
-
- Safety and efficient allocation is hard
 - Several simplifications here, e.g., only one unit of resource requested at a time
 - No focus on efficiency, resource is modeled as a set

Next slides: two versions

1. Allocator as (passive) monitor
2. Allocator as active monitor

Recap: Semaphore Monitor Passing the Condition

Await

```
monitor Semaphore { # monitor invariant: s >= 0
    int s := 0;           # value of the semaphore
    cond pos;            # wait condition

    procedure Psem() {
        if (s=0) wait(pos);
        else      s := s - 1; }

    procedure Vsem() {
        if (empty(pos)) s := s + 1;
        else              signal(pos); }

}
```

Allocator as a (Passive) Monitor

Await

```
monitor Resource_Allocator {
    int avail := MAXUNITS;
    set units;
    cond free;           // signalled when process wants a unit

    procedure acquire(int &id) {
        if (avail = 0) wait(free);
        else           avail := avail -1;
        remove(units, id); } // exact management abstracted here

    procedure release(int id) {
        insert(units, id);
        if (empty(free)) avail := avail +1;
        else             signal(free); }

}
```

Allocator as a Server Process: Code-Design Process for Monitors

1. Interface and internal variables
 - 1.1 Two types of operations: get unit, free unit
 - 1.2 One request channel *encoded* in the arguments to a request.
2. Control structure
 - 2.1 First check the kind of requested operation,
 - 2.2 Then, perform resource management for that operation
3. Synchronization, scheduling, and mutex
 - 3.1 Cannot wait (ie. `wait(free)`) when no unit is free.
 - 3.2 Must save the request and return to it later
 - ⇒ queue of pending requests (**queue; insert, remove**).
 - 3.3 Upon request: synchronous/blocking call ⇒ “ack”-message back
 - 3.4 No internal parallelism due to mutex

Channel Declarations

Await

```
type op_kind = enum(ACQUIRE, RELEASE);
chan request(int clientID, op_kind kind, int unitID);
chan reply[n](int unitID);

process Client[i = 0 to n-1] {
    int unitID;
    send request(i, ACQUIRE, 0); // make request
    receive reply[i](unitID);      // works as ``if synchronous''
    ...                           // use resource unitID
    send request(i, RELEASE, unitID); // free resource
    ...
}
```

Note the problems with type-uniform channels: ACQUIRE request does not use its last parameter, RELEASE does not use the first one.

Await

```
process Resource_Allocator {
    int avail := MAXUNITS;
    set units := ...;           // initial value
    queue pending;              // initially empty
    int clientID, unitID; op_kind kind; ...
    while (true) {
        receive request(clientID, kind, unitID);
        if (kind = ACQUIRE) {
            if (avail = 0) insert(pending, clientID); // save request
            else { // perform request now
                avail:= avail-1;
                remove(units, unitID);
                send reply[clientID](unitID); } }
        else { // kind = RELEASE
            if empty(pending) avail := avail+1; insert(units, unitID);
            else { // allocates to waiting client
                remove(pending, clientID);
                send reply[clientID](unitID); } } } }
```

Duality: Mainonitors & Message Passing

monitor-based programs	message-based programs
monitor variables	local server variables
process-IDs	request channel, operation types
procedure call	send request(), receive reply[i]()
go into a monitor	receive request()
procedure return	send reply[i]()
wait statement	save pending requests in a queue
signal statement	get and process pending request (reply)
procedure body	branches in branching over op. type

Synchronous Message Passing

Synchronous Channels

- Asynchronous channels pass messages, but do not synchronize two processes
- Next: Synchronous channels
- Natural connection to barriers

Primitives

```
synch_send c(expr1, ..., exprN);
```

- Sender waits until message is received via the channel,
- Sender and receiver synchronize by the sending and receiving of message
- Same receiving primitive

Synchronous Message Passing: Discussion

Advantages

- Gives maximum *size* of channel (for fixed number of processes), as sender synchronizes with receiver
 - Receiver has at most 1 pending message per channel per sender
 - Each sender has at most 1 unsent message

Disadvantages

- Reduced parallelism: when 2 processes communicate, 1 is always blocked
- Higher risk of *deadlock*

Example: Blocking with Synchronous Message Passing

Await

```
chan values(int);  
  
process Producer {  
    int data[n];  
    for (i = 0 to n-1) {  
        ... //computation  
        synch_send values(data[i]); }  
}  
  
process Consumer {  
    int results[n];  
    for (i = 0 to n-1) {  
        receive values(results[i]);  
        ... //computation  
    } }
```

- Assume both producer and consumer vary in time complexity.
- Communication using `synch_send/receive` will **block**.
- With *asynchronous* message passing, the waiting is reduced.

Example: Deadlock using Synchronous Message Passing

Await

```
chan in1(int), in2(int);

process P1 {
    int v1 = 1, v2;
    synch_send in2(v1);
    receive in1(v2);}

process P2 {
    int v1, v2 = 2;
    synch_send in1(v2);
    receive in2(v1);}
```

- P1 and P2 both block on `synch_send` – program *deadlocks*
- One process must be modified to do `receive` first ⇒ asymmetric solution.
- With asynchronous channels, all goes well

Encoding

- Despite all, many implementations (e.g., Go) and theories (e.g., π -calculus have *synchronous channels*)
- Main reason: It is easier to encode asynchronous message passing with synchronous channels than vice versa
- Requires way to spawn new thread/process

Await

```
chan v(int);  
  
process Send{  
    spawn { synch_send v(1); } //spawns new thread and continues  
}  
process Receive {  
    int res;  
    receive v(res);  
}
```

Summary

Today's lecture

- Shared memory vs. distributed memory
- Synchronous and asynchronous message passing, the high level picture
- *Asynchronous message passing*: channels, messages, primitives
- Example: filters and sorting networks
- Comparison of message passing and monitors
- Basics *synchronous message passing*

Next lectures in this module

- Concurrency in Go
- Actors with asynchronous communication / Await primitive