

# **Actors, Active Objects and Asynchronous Communication**

---

Andrea Pferscher

14.10.2024

University of Oslo

## Part 2: Message Passing

### Structure

- Part 1: Shared Memory (and Java)
- **Part 2: Message Passing (and Go)**
- Part 3: Analyses and Tool Support (and maybe Rust)

### Content of this part:

- Synchronous and asynchronous message passing
- Channels, actors, go-routines, asynchronous programming

### Outline Today

- Actors
- Futures and promises
- Active objects
- Asynchronous communication with await-statement

## Message Passing and Channels

- Shared memory vs. distributed memory
- Synchronous and asynchronous message passing, the high level picture
- *Asynchronous message passing*: channels, messages, primitives
- Example: filters and sorting networks
- Comparison of message passing and monitors
- Basics *synchronous message passing*

## Actors

---

# Async. Communication without Channels

## Channels

- Need additional primitives for concurrency; send and receive
- Channels are explicit while process/objects are implicit
- Complex typing disciplines

**Can we do asynchronous communication without explicit channels?**

- Actors: Messages between objects
- Active Objects: Messages between objects with cooperative scheduling
- Async/Await in mainstream languages: Using (lightweight) threads (with shared memory)

# Actors

- Actors: a programming concept for distributed concurrency which combines a number of topics we have discussed in the course;
  - active monitors,
  - objects and encapsulation,
  - race-free (no race conditions on shared state)
- Examples of programming languages that implement actors:  
Erlang, Scala's Akka library, Dart, Swift, etc.

# Object-Oriented Programming and Language Design

## What are objects

How do OO programs fit into the design of programming languages?

- **State space:** local or global?
- **Thread interaction and objects**
- **Communication:** shared variables, channels or messages?
- **Communication:** synchronous or asynchronous?
- **Dynamic state allocation:** object creation

## What can we do to protect objects against races?

## Can we combine objects with ideas from monitors?

- Passive monitors vs. active monitors
- A method is *active*, if a statement in the method is executed by some thread

# Passive Monitors – Repetition

## Await

```
monitor name {  
    monitor variables  
    ## monitor invariant  
    initialization code  
    procedures  
}
```

- Threads *communicate* by calling monitor methods
- Threads do not need to know all the implementation details: only the procedure names are visible from outside the monitor
- Statements *inside* a monitor: *no* access to variables *outside* the monitor
- Statements *outside* a monitor: *no* access to variables *inside* the monitor
- **Monitor variables:** *initialized* before the monitor is used
- **Monitor invariant:** describes a condition on the inner state
- The monitor invariant can be analyzed by sequential reasoning inside the monitor

## Passive Monitors: Synchronization with condition variables – Repetition

- Monitors contain *special* type of variables: **cond** (condition)
- Used for synchronization/to delay processes
- Each such variable is associated with a *wait condition*
- The value of a condition variable: *queue* of delayed threads
- Not directly accessible by programmer, instead, manipulated by special operations

```
cond cv;          # declares a condition variable cv
empty(cv);       # asks if the queue on cv is empty
wait(cv);        # causes thread to wait in the cv queue
signal(cv);      # wakes up a thread in the queue to cv
signal_all(cv); # wakes up all threads in the cv queue
```

# Passive Monitors – Repetition

## Await

```
monitor Mon { // monitor invariant: r ≥ 0
    int r := 0 // number of resources
    cond res; // wait condition variable

    procedure Acquire() {
        while (r=0) { wait (res); }
        r := r - 1 }

    procedure Release() {
        r := r+1;
        signal (res); }}
```

- wait and signal: *FIFO signaling strategy*
- A thread in the monitor can execute signal(cv).  
If there is a waiting thread, do we get *two active methods* in the monitor?

# Objects as Passive Monitors in Java

Java

```
class Mon { // class invariant: this.r >= 0
    int r = 0; // number of resources
    Condition res; // wait condition
    public synchronized void Acquire() {
        while( r == 0 ) { res.await(); };
        r = r - 1;
    }
    public synchronized void Release() {
        r = r + 1;
        res.signal();
    }
}
```

- How do condition variables and synchronized methods relate?

# Actors

**Fundamental idea:** Decouple communication and control.

## Capabilities of Actors

An actor reacts to incoming messages to

- change its state,
- send a finite number of messages to other actors, and
- create a finite number of new actors.

## Intuition

We can think of an actor as an object that can only communicate asynchronously.

Some actor models can also pattern match over its message queue of incoming messages.

# Implementation of Actors in Programming Languages

- Supported by numerous languages and frameworks
  - Not always strictly OO: Erlang, ...
  - Sometimes as library, not part of language: Akka actors, ...
  - Numerous differences on how basic capabilities are implemented or extended
- 
- Type safety: Can we guarantee statically whether messages can be processed?
  - Integration with OO: Are messages methods? Do actors have a class?
  - Integration with other primitives: Can actors share state?
  - Integration with error handling: What happens when an actor fails?
  - Here: foundations

# Actors: Communication & Concurrency

## Actors

- Recipients of messages are identified by name (no channels).
- An actor can only communicate with actors that it knows.
- An actor can obtain names from messages that it receives, or because it has created the actor

The actor model is characterized by

- inherent concurrency among actors
- dynamic creation of actors,
- inclusion of actor names in messages, and
- interaction only through direct asynchronous message passing with no restriction on message arrival order.
- *message servers* might be implemented by matching messages from the queue to procedures

## Example: Erlang-style Actors - Matching Messages

### Publish and Subscribe Server

```
runServer(Subs) ->
    receive
        {sub,from} -> runServer(Subs + from); % subscribe
        {publish,value} -> % publish
            for(id in Subs) id!{value}, % broadcast value
            runServer(Subs);
        _ -> runServer(Subs); % ignore other messages
    end.
```

```
Server { % publish and subscribe server
    start() -> spawn(fun() -> runServer([])).} % start the server
```

```
Client { % send requests to the server
    start() -> Server!{sub,self}, Server!{publish,10}.}
```

## Example: Erlang-style Actors

- State as argument to recursive calls
- We can dynamically change the message server
- An actor can match different messages in different states
- ... but tricky to detect errors in message servers

```
runServer1(Subs) -> receive % subscribe when there is space
    {sub,from} -> if(size(Subs) >= 9) runServer2(Subs + from)
                    else runServer1(Subs + from);
    {unsub,from} -> runServer1(Subs - from);
    ...

```

## Example: Erlang-style Actors – Handling Return Values between Actors

```
id1 = spawn(fun() -> func1([])); id2 = spawn(fun() -> func2([]))
id1!{step1, 42, id2};

...
func1(history) -> receive
    {step1, data, other} -> newData = doSomethingFirst(data),
                                other!{step2, newData, self},
                                func1(insert(history,data));
    {step3, data, other} -> newData = doSomethingThird(data),
                                other!{step4, newData, self},
                                func1(insert(history,data));

func2(history) -> receive
    {step2, data, other} -> newData = doSomethingSecond(data),
                                other!{step3, newData, self},
                                func2(insert(history,data)); ...
```

## Futures

---

# Futures – Handling Return Values between Actors

## Welcome to “callback hell”!

- Problem: Logically related code is scattered in program
- We need a way to identify callback messages
- We also need a way to wait for a result
- Solution: futures, special mailboxes transmit return values

Reminder in Java:

Java

```
ExecutorService service = Executors.newFixedThreadPool(2);
Future<Int> f = service.submit(() -> { /* do */ return 1;});
...
Int x = f.get(); //essentially a join
```

# Futures and Promises

## Futures

- It is a handle for the caller of a process. It will contain the result value once computed
- It can be read multiple times
- It can be used by the caller to synchronize with the callee

### Java

```
Future<Int> f = service.submit(() -> { return 1;});  
...  
Int x = f.get();
```

## Promises

- What if the value will be computed somewhere else?
- A *promise* is a future for which it is not clear who computes the value

# Promises

A promise:

- May be eventually completed (but maybe by somebody else)
- Must be completed (written) only once
- Deadlock/starvation occurs if it is never completed
- It can be seen as a handle for the callee and the callee does not synchronise with the caller

Java calls promises *CompletableFuture*:

Java

```
CompletableFuture<Integer> f = new CompletableFuture<>();
service.submit(() -> { f.complete(1); return null;});
...
Int x = f.get();
```

## Promises – Example: Service Delegation

Java

```
/* the function casts a promise as a future */
/* from outside the future can only be retrieved */
Future<Integer> callAsync() ... {
    CompletableFuture<Integer> completableFuture = new CompletableFuture<>();
    service1.submit(() -> {
        if(/* service1 cannot process, then it delegates to service2 */)
            service2.submit(() ->
                { /* compute */ completableFuture.complete(1); return null })
        else { /* process the request */
            /* compute */ completableFuture.complete(1); }
        return null;
    });
    return completableFuture;
}
```

# Composition Futures/Promises

Logically related Futures/Promises scattered in the code.

Java

```
CompletableFuture<Integer> f1 = CompletableFuture.supplyAsync(() -> 1);
...
CompletableFuture<Integer> f2
    = CompletableFuture.supplyAsync(() -> f1.get + 1);
```

Connecting Futures/Promises (composition)

Java

```
CompletableFuture<Integer> f
    = CompletableFuture.supplyAsync(() -> 1)
        .thenApply((res) -> res + 1);
```

Very similar patterns are common in web development with JavaScript

# Interpreting Futures/Promises as Channels

## Channel-view on single-read futures

- Create channel and send it via an asynchronous message
- For the caller, the channel behaves as a future:  
caller waits on the channel for a return (caller side does not write on the channel).
- For the callee, the channel behaves as a promise:  
it can be passed around, and eventually someone will write on it *exactly once*  
(callee side does not read on the channel)

## Limits of this view

- Futures may be read more than once
- “immediately creating and sharing a channel” may be more complex and its implementation is delegated to the programmer

## Active Objects

---

# Motivation

- How to combine monitors and actors?
- How to make signalling less error-prone?
- How to make conditions/invariants easier to use?
- How to connect futures/promises with actors?

## Active Objects

An active object<sup>a</sup> is an actor with an *implicit* message server, that only communicates asynchronously, but allows internal message handlers to use *cooperative scheduling*.

- One process/thread per object
- Messages identified with methods
- Implicit queue of tasks (procedures in the methods)
- Explicit synchronization

---

<sup>a</sup>ABS is a modelling language to run simulations of distributed systems.

The simulation tool is maintained by the PSY group at IFI: <https://abs-models.org/>

# Cooperative concurrency

## Active Monitors as Active Objects

- Cooperative concurrency:  
constructs to suspend and resume execution (=task) of a local method
- External cooperation (operations on futures)
  - Send is **asynchronous**: `Fut < T > f = o!m(...); ... ;`
  - Retrieve value is **blocking**: `x = f.get;`
  - Check for value is **suspending**: `await f?`
- Interaction patterns between methods
  - `Fut < T > f = o!m(...);x = f.get;`
  - `Fut < T > f = o!m(...); ...; x = f.get;`
  - `Fut < T > f = o!m(...); ...; await f?; x = f.get;`

## Cooperative Scheduling – Example: The Diner

- Each object runs one thread and each method call spawns a *task*
- Thread is responsible to schedule tasks in some order
- Waiting on future suspends the task, not the thread!
- Reading on future potentially blocks task and thread – no other task can run

ABS

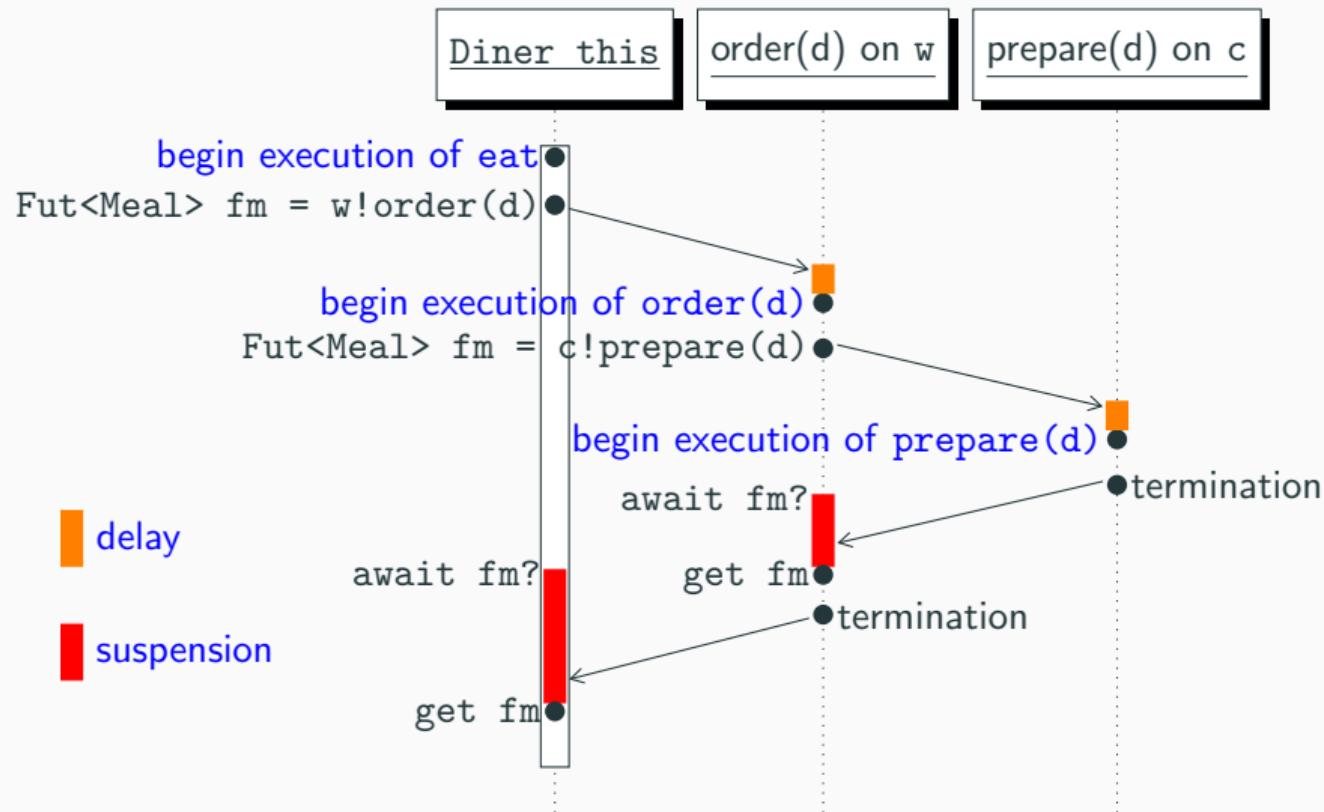
```
class Diner(IWaiter w) implements IDiner {  
    Unit eat(Dish d) {  
        Fut<Meal> fm = w!order(d); // place order with waiter  
        await fm?; // while waiting do something else, e.g., take a phone call  
        Meal m = fm.get; // receive meal  
        Fut<Unit> fc = this!consume(m);  
        Fut<Unit> fp = w!pay(this, d); // eating, paying in some order  
        await fc? & fp?; // eaten and paid – ready to leave!  
    }  
    Unit takeCall(){ ... }  
    Int takeMoney(Int a) { ... }  
    ...  
}
```

## Example (Continuation): – The Waiter

ABS

```
class Waiter(ICook c, Int purse) implements IWaiter {
    Meal order(Dish d) {
        Fut<Meal> fm = c!prepare(d); // place order with cook
        await fm?; // waiter serves other guests while meal is cooked
        Meal m = fm.get(); // receive meal reuse names for local variables
        return m; // ready to serve the meal!
    }
    Unit pay(IDiner g, Dish d) {
        Int amount = price(d); // lookup price in the menu
        Int a = g.takeMoney(amount); // synchronous (blocking) call, no wait
        this.purse = this.purse + amount; // no data race possible
    }
}
```

## Example (Continuation) – The Restaurant Experience



# Condition Synchronization

- Condition variables can be derived from monitor invariant
- Or can be bound to some other condition
- Error-prone implementations
- Active Object approach: condition synchronization as primitive

ABS

```
class C () {  
    int i = 0;  
    Unit inc() { i = i+1; return; }  
    Int isGreaterThanTen(){ await i > 10; return i; }  
}
```

- Condition variables: explicit suspension instead of busy waiting
- Every time the object is idle, the object thread evaluates all conditions of suspended tasks, otherwise it waits for new messages to arrive

## Objects as Passive Monitors (reminder) – Example

Java

```
class Mon {  
    int r = 0;  
    Condition res;  
  
    public synchronized void Acquire() {  
        while( r == 0 ) { res.await(); }  
        r = r - 1;  
    }  
  
    public synchronized void Release() {  
        r = r + 1;  
        res.signal();  
    }  
}
```

## Monitors with active objects – Example

ABS

```
class Mon {  
    int r = 0  
  
    Unit Acquire() {  
        await (r!=0);  
        r = r - 1;  
    }  
  
    Unit Release() {  
        r = r+1;  
    }  
}
```

- With cooperative concurrency, we can avoid error-prone signaling in the monitor.
- The active object only has one queue, but reactivation of Acquire methods can only happen when the await-condition holds

# Bounded Buffer Synchronization with Active Objects(1)

Let us now solve the bounded buffer problem with active objects

## Bounded buffer synchronization

- buffer of size  $n$  (“channel”, “pipe”)
- producer: performs put operations on the buffer.
- consumer: performs getVal operations on the buffer.
- two access operations (“methods”)
  - put operations must wait if buffer full
  - getVal operations must wait if buffer empty

## Bounded Buffer Synchronization with Active Objects (2)

ABS

```
class Bounded Buffer (Int n) {
    List<T> buf = [];
    Unit put(T data){
        await (length(buf) < n);
        buf = appendright(buf,data);
    }
    T getVal() {
        await (length(buf) > 0);
        T tmp = head(buf); buf = tail(buf); return tmp;
    }
}
```

## What is a deadlock?

A system is deadlocked if it is *stuck*:  
It cannot continue execution, and  
it has not finished its execution.

A system is deadlocked if there is a circular dependency: There is a sequence of components  $C_1, \dots, C_n$ , such that  $C_i$  depends on  $C_{i+1}$  before it can continue and  $C_n$  depends on  $C_1$ .

- Actors without futures/channels cannot deadlock – they can always continue execution  
... but there can be messages that cannot be processed with the current message server!
- In some concurrency models, a system can only get stuck because of a circular dependency

## Local Dependencies – Between the Object and its Tasks

ABS

```
class C (){

    Unit m(){
        Fut<T> f = this!n();
        f.get; // deadlock
    }

    T n(){ /* do some computation */ return value; }

}
```

# Dependencies due to Synchronization Between Tasks

- A task depends on another task if it waits for its future

ABS

```
class C {  
    Fut<Unit> f1;  
    Unit store(Fut<Unit> fut) { f1 = fut; }  
    Unit m(){ await f1?; return; }} //depends on d.n  
  
class D(C c) {  
    Unit n(){ Fut<Unit> f2= c!m();  
              await f2?; //depends on c.m  
              return; } }  
{ // Main block  
    C c = new C(); D d = new D(c);  
    Fut<Unit> f;  
    await c!store(f); f= d!n(); // deadlock  
}
```

## Dependencies Related to the State of an Object

- In a given state a task  $t_1$ , that might be stuck on condition  $e_1$ , depends on another task  $t_2$ , that might be stuck on condition  $e_2$ .
- Here  $e_1$  and  $e_2$  are conditions related to the state of an object, which create dependencies between the tasks.

ABS

```
class D { // here exclamation mark is negation
    Bool b1 = false; Bool b2 = false;
    Unit m(){ b1 = true; await b2; b1 = !b2;}
    Unit n(){ await !b1; b2 = !b1;}
}
```

- What happens if we call  $n()$  and then  $m()$  on a  $D$ -object?
- There is no procedure to decide whether an arbitrary program ever deadlocks because it depends on the scheduling of tasks

# Outlook: Analysis and Modelling

## Reasoning

Monitors and actors are well-suited for manual and automatic reasoning

- Built-in mutex ensures that between interaction points, code can be seen as sequential
- Sequential reasoning has to be extended only at these points
- Full concurrency requires non-local reasoning at every point

## Programming is Modelling

A program can be used to model a part of the world.

- A program analysis then can be used to derive properties over the world
- For example, 5 philosophers programs are *executable* models
- Allows analysis for deadlock freedom.

## Async/Await

---

# Recap on Message Passing

## Message Passing So Far

- Channels: Asynchronous shared entities
- Actors: Monitors that send asynchronous messages
- Active Objects: Monitors with their own thread that send asynchronous messages

# Java and Async/Await

Reminder in Java:

Java

```
ExecutorService service = Executors.newFixedThreadPool(2);
Future<Int> f = service.submit(() -> { /* do */ return 1;});
...
Int x = f.get(); //essentially a join
```

- Executed function disconnected from classes
- Much boilerplate code, especially when call-backs are involved
- Asynchronous code (library) does not mirror synchronous code (language constructs)

# C# and Async/Await

## C#'s Asynchronous Concurrency

- Better abstraction to handle Futures/Tasks.
  - Concurrency as first-class construct of language
- 
- Methods annotated with `async` can only be called asynchronously
  - Methods annotated with `async` return a `Task`
  - Only methods annotated with `async` can perform an `await`
  - Expression `await` suspends the thread until the task has finished.

## C# and Async/Await – Example: Comp. Two Numbers

- Example: Reading two numbers from user and performing some long-lasting computation
- Synchronous version
- Await version: Note that Method must be async to use await
- Asynchronous version: Now both reads can be concurrent

C#

```
class C{
    void async Method() {
        Task<Int> t1 = GetFirstNumber(); Task<Int> t2 = GetSecondNumber();
        Int i1 = await t1; Int i2 = await t2;
        Int res = await Compute(i1,i2);
    }
    Task<Int> async GetFirstNumber() {...}
    ...
}
```

# Pros and Cons of Async/Await

## What Color is your Function

- Only async methods can access results of async methods
- Separates whole program into two sets of methods that can only interact at specific points
- Sometimes called colored-function problem, after a popular blog entry<sup>a</sup>

---

<sup>a</sup><https://journal.stuffwithstuff.com/2015/02/01/what-color-is-your-function/>

- Forces programmer to think about concurrency
- Can still use threads and tasks directly to circumvent all this

# Wrap-Up

## Today's Lecture

- Actors – Monitors with message passing
- Futures/Promises – Handling asynchronous results
- Active Objects – Actors with cooperative concurrency and futures
- Async/Await – Language-integrated asynchronicity with threads and futures

## Next Lectures

- Next Block: How to type channels?

Note: ABS example courtesy of Reiner Hähnle